

# akulas\_HW2

January 16, 2019

- Adam Kulas
- watiam: akulas
- student number: 20302000

References: [About Feature Scaling and Normalization](#), Sebastian Raschka

```
In [1]: import pandas as pd
import numpy as np
from sklearn import preprocessing
from IPython.display import display, HTML
```

```
In [2]: data = pd.read_csv('data/winequality-red.csv', sep=';')
data = data.iloc[:10, :]
display(data)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
5	7.4	0.66	0.00	1.8	0.075	
6	7.9	0.60	0.06	1.6	0.069	
7	7.3	0.65	0.00	1.2	0.065	
8	7.8	0.58	0.02	2.0	0.073	
9	7.5	0.50	0.36	6.1	0.071	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	
5	13.0	40.0	0.9978	3.51	0.56	
6	15.0	59.0	0.9964	3.30	0.46	
7	15.0	21.0	0.9946	3.39	0.47	
8	9.0	18.0	0.9968	3.36	0.57	
9	17.0	102.0	0.9978	3.35	0.80	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
5	9.4	5
6	9.4	5
7	10.0	7
8	9.5	7
9	10.5	5

## 1 Calculate the Min Max Scaling of the Data

```
In [3]: # Calculating the min max scaling manually
        datamm = (data - data.min()) / (data.max() - data.min())
```

```
In [4]: #Calculating min max scaling using sklearn
        minmax_scale = preprocessing.MinMaxScaler().fit(data)
        data_minmax = minmax_scale.transform(data)
        data_minmax = pd.DataFrame(data_minmax)
        data_minmax.columns = data.columns
        display(data_minmax)
```

```
/home/calkulas/anaconda3/envs/ece657A/lib/python3.7/site-packages/sklearn/preprocessing/data.py
    return self.partial_fit(X, y)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides \
0	0.025641	0.700000	0.000000	0.142857	0.333333
1	0.128205	1.000000	0.000000	0.285714	1.000000
2	0.128205	0.800000	0.071429	0.224490	0.818182
3	1.000000	0.000000	1.000000	0.142857	0.303030
4	0.025641	0.700000	0.000000	0.142857	0.333333
5	0.025641	0.633333	0.000000	0.122449	0.303030
6	0.153846	0.533333	0.107143	0.081633	0.121212
7	0.000000	0.616667	0.000000	0.000000	0.000000
8	0.128205	0.500000	0.035714	0.163265	0.242424
9	0.051282	0.366667	0.642857	1.000000	0.181818

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	0.125	0.190476	0.941176	1.000000	0.294118
1	1.000	0.583333	0.647059	0.114286	0.647059
2	0.375	0.428571	0.705882	0.285714	0.558824
3	0.500	0.500000	1.000000	0.000000	0.352941
4	0.125	0.190476	0.941176	1.000000	0.294118

5	0.250	0.261905	0.941176	1.000000	0.294118
6	0.375	0.488095	0.529412	0.400000	0.000000
7	0.375	0.035714	0.000000	0.657143	0.029412
8	0.000	0.000000	0.647059	0.571429	0.323529
9	0.500	1.000000	0.941176	0.542857	1.000000

	alcohol	quality
0	0.000000	0.0
1	0.363636	0.0
2	0.363636	0.0
3	0.363636	0.5
4	0.000000	0.0
5	0.000000	0.0
6	0.000000	0.0
7	0.545455	1.0
8	0.090909	1.0
9	1.000000	0.0

## 2 Calculate the Z score of the data

```
In [5]: # calculating the z score manually
        # print(data.mean(), data.std(), data)
        data_zscore = (data - data.mean()) / data.std()
        #display(data_zscore)
```

```
In [6]: std_scale = preprocessing.StandardScaler().fit(data)
        data_zscore = pd.DataFrame(std_scale.transform(data))
        data_zscore.columns = data.columns
        display(data_zscore)
```

```
/home/calkulas/anaconda3/envs/ece657A/lib/python3.7/site-packages/sklearn/preprocessing/data.py
    return self.partial_fit(X, y)
/home/calkulas/anaconda3/envs/ece657A/lib/python3.7/site-packages/ipykernel_launcher.py:2: Data
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	-0.498662	0.451753	-0.563489	-0.329398	-0.103362	
1	-0.135999	1.630239	-0.563489	0.206831	2.170608	
2	-0.135999	0.844582	-0.346763	-0.022981	1.550434	
3	2.946642	-2.298048	2.470683	-0.329398	-0.206725	
4	-0.498662	0.451753	-0.563489	-0.329398	-0.103362	
5	-0.498662	0.189867	-0.563489	-0.406002	-0.206725	
6	-0.045333	-0.202961	-0.238399	-0.559211	-0.826898	
7	-0.589328	0.124396	-0.563489	-0.865627	-1.240347	
8	-0.135999	-0.333904	-0.455126	-0.252794	-0.413449	
9	-0.407997	-0.857676	1.387050	2.887978	-0.620174	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	-0.896665	-0.626576	0.731200	1.276466	-0.305196	
1	2.406839	0.761143	-0.284356	-1.276466	0.957683	
2	0.047193	0.214466	-0.081244	-0.782350	0.641963	
3	0.519122	0.466778	0.934311	-1.605877	-0.094716	
4	-0.896665	-0.626576	0.731200	1.276466	-0.305196	
5	-0.424736	-0.374264	0.731200	1.276466	-0.305196	
6	0.047193	0.424726	-0.690578	-0.452940	-1.357594	
7	0.047193	-1.173253	-2.518578	0.288234	-1.252354	
8	-1.368595	-1.299409	-0.284356	0.041176	-0.199956	
9	0.519122	2.232966	0.731200	-0.041176	2.220561	

	alcohol	quality
0	-0.880830	-0.620174
1	0.293610	-0.620174
2	0.293610	-0.620174
3	0.293610	0.620174
4	-0.880830	-0.620174
5	-0.880830	-0.620174
6	-0.880830	-0.620174
7	0.880830	1.860521
8	-0.587220	1.860521
9	2.348881	-0.620174

### 3 Calculate the mean subtracted normalized values

```
In [7]: data_meansub = data - data.mean()
display(data_meansub)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	-0.55	0.069	-0.104	-0.43	-0.001	
1	-0.15	0.249	-0.104	0.27	0.021	
2	-0.15	0.129	-0.064	-0.03	0.015	
3	3.25	-0.351	0.456	-0.43	-0.002	
4	-0.55	0.069	-0.104	-0.43	-0.001	
5	-0.55	0.029	-0.104	-0.53	-0.002	
6	-0.05	-0.031	-0.044	-0.73	-0.008	
7	-0.65	0.019	-0.104	-1.13	-0.012	
8	-0.15	-0.051	-0.084	-0.33	-0.004	
9	-0.45	-0.131	0.256	3.77	-0.006	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	-3.8	-14.9	0.00072	0.155	-0.029	
1	10.2	18.1	-0.00028	-0.155	0.091	
2	0.2	5.1	-0.00008	-0.095	0.061	

3	2.2	11.1	0.00092	-0.195	-0.009
4	-3.8	-14.9	0.00072	0.155	-0.029
5	-1.8	-8.9	0.00072	0.155	-0.029
6	0.2	10.1	-0.00068	-0.055	-0.129
7	0.2	-27.9	-0.00248	0.035	-0.119
8	-5.8	-30.9	-0.00028	0.005	-0.019
9	2.2	53.1	0.00072	-0.005	0.211

	alcohol	quality
0	-0.3	-0.5
1	0.1	-0.5
2	0.1	-0.5
3	0.1	0.5
4	-0.3	-0.5
5	-0.3	-0.5
6	-0.3	-0.5
7	0.3	1.5
8	-0.2	1.5
9	0.8	-0.5

## 4 Calcualte distance of data

In [8]: `from scipy.spatial import distance`

### 4.1 Manhattan Distance

In [9]: `# calculate the pairwise distance using the manhattan distance metric`  
`d = pd.Series(distance.pdist(data, 'cityblock'))`  
`dsq = pd.DataFrame(distance.squareform(d))`  
`np.fill_diagonal(dsq.values, np.nan)`  
`display(dsq)`

	0	1	2	3	4	5	6	7	\
0	NaN	49.1330	25.6568	38.5512	0.0000	8.1410	30.2784	20.6742	
1	49.1330	NaN	23.5562	21.4242	49.1330	41.2740	20.1894	60.7652	
2	25.6568	23.5562	NaN	13.9880	25.6568	17.7978	6.6336	37.2894	
3	38.5512	21.4242	13.9880	NaN	38.5512	30.6102	9.0876	48.0834	
4	0.0000	49.1330	25.6568	38.5512	NaN	8.1410	30.2784	20.6742	
5	8.1410	41.2740	17.7978	30.6102	8.1410	NaN	22.1374	24.5332	
6	30.2784	20.1894	6.6336	9.0876	30.2784	22.1374	NaN	41.8158	
7	20.6742	60.7652	37.2894	48.0834	20.6742	24.5332	41.8158	NaN	
8	20.9040	68.5150	44.9992	55.8532	20.9040	28.9630	49.8344	11.0302	
9	80.3650	48.5380	55.6418	52.4342	80.3650	72.4240	51.7934	91.4892	

  

	8	9
0	20.9040	80.3650
1	68.5150	48.5380

```

2  44.9992  55.6418
3  55.8532  52.4342
4  20.9040  80.3650
5  28.9630  72.4240
6  49.8344  51.7934
7  11.0302  91.4892
8      NaN 100.0630
9 100.0630      NaN

```

```

In [10]: distances = pd.DataFrame([dsq.idxmin(axis=0), dsq.min(axis=0), dsq.idxmax(axis=0), dsq.max(axis=0)],
    distances.columns = ['nearest_manhattan_point', 'near_manhattan_dist_val', 'far_manhattan_point', 'far_manhattan_dist_val']
    display(distances)

```

	nearest_manhattan_point	near_manhattan_dist_val	far_manhattan_point	\
0	4.0	0.0000	9.0	
1	6.0	20.1894	8.0	
2	6.0	6.6336	9.0	
3	6.0	9.0876	8.0	
4	0.0	0.0000	9.0	
5	0.0	8.1410	9.0	
6	2.0	6.6336	9.0	
7	8.0	11.0302	9.0	
8	7.0	11.0302	9.0	
9	1.0	48.5380	8.0	

	far_manhattan_dist_val
0	80.3650
1	68.5150
2	55.6418
3	55.8532
4	80.3650
5	72.4240
6	51.7934
7	91.4892
8	100.0630
9	100.0630

## 4.2 Cosine Distance

```

In [11]: # calculate the pairwise distance using the cosine distance metric
    d = pd.Series(distance.pdist(data, 'cosine'))
    dsq = pd.DataFrame(distance.squareform(d))
    np.fill_diagonal(dsq.values, np.nan)
    display(dsq)

```

	0	1	2	3	4	5	6	\
0	NaN	0.015207	0.007749	0.007776	0.000000	0.001322	0.011912	

1	0.015207	NaN	0.004858	0.005985	0.015207	0.007860	0.006094
2	0.007749	0.004858	NaN	0.001097	0.007749	0.003120	0.000601
3	0.007776	0.005985	0.001097	NaN	0.007776	0.003502	0.001635
4	0.000000	0.015207	0.007749	0.007776	NaN	0.001322	0.011912
5	0.001322	0.007860	0.003120	0.003502	0.001322	NaN	0.006079
6	0.011912	0.006094	0.000601	0.001635	0.011912	0.006079	NaN
7	0.051499	0.082365	0.089001	0.086412	0.051499	0.060757	0.102130
8	0.045682	0.099274	0.088973	0.085518	0.045682	0.060301	0.101541
9	0.035050	0.019639	0.011228	0.014255	0.035050	0.025574	0.007549

	7	8	9
0	0.051499	0.045682	0.035050
1	0.082365	0.099274	0.019639
2	0.089001	0.088973	0.011228
3	0.086412	0.085518	0.014255
4	0.051499	0.045682	0.035050
5	0.060757	0.060301	0.025574
6	0.102130	0.101541	0.007549
7	NaN	0.015776	0.159302
8	0.015776	NaN	0.153601
9	0.159302	0.153601	NaN

```
In [12]: distances = pd.DataFrame([dsq.idxmin(axis=0), dsq.min(axis=0), dsq.idxmax(axis=0), dsq.max(axis=0)],
    distances.columns = ['nearest_cosine_point', 'near_cosine_dist_val', 'far_cosine_point', 'far_cosine_dist_val']
    display(distances)
```

	nearest_cosine_point	near_cosine_dist_val	far_cosine_point	\
0	4.0	0.000000	7.0	
1	2.0	0.004858	8.0	
2	6.0	0.000601	7.0	
3	2.0	0.001097	7.0	
4	0.0	0.000000	7.0	
5	0.0	0.001322	7.0	
6	2.0	0.000601	7.0	
7	8.0	0.015776	9.0	
8	7.0	0.015776	9.0	
9	6.0	0.007549	7.0	

	far_cosine_dist_val
0	0.051499
1	0.099274
2	0.089001
3	0.086412
4	0.051499
5	0.060757
6	0.102130
7	0.159302

```

8          0.153601
9          0.159302

```

### 4.3 Euclidean Distance

```

In [13]: # calculate the pairwise distance using the euclidean distance metric
         d = pd.Series(distance.pdist(data, 'euclidean'))
         dsq = pd.DataFrame(distance.squareform(d))
         np.fill_diagonal(dsq.values, np.nan)
         display(dsq)

```

	0	1	2	3	4	5 \
0	NaN	35.860192	20.409705	26.985420	0.000000	6.325472
1	35.860192	NaN	16.404589	11.257696	35.860192	29.565511
2	20.409705	16.404589	NaN	7.296300	20.409705	14.165186
3	26.985420	11.257696	7.296300	NaN	26.985420	20.789202
4	0.000000	35.860192	20.409705	26.985420	NaN	6.325472
5	6.325472	29.565511	14.165186	20.789202	6.325472	NaN
6	25.326029	12.857342	5.071906	4.186459	25.326029	19.114166
7	13.779881	47.142170	33.084192	39.271562	13.779881	19.228955
8	16.254766	51.590491	36.554474	42.907828	16.254766	22.455276
9	68.404041	36.085200	48.199803	42.390953	68.404041	62.289232

  

	6	7	8	9
0	25.326029	13.779881	16.254766	68.404041
1	12.857342	47.142170	51.590491	36.085200
2	5.071906	33.084192	36.554474	48.199803
3	4.186459	39.271562	42.907828	42.390953
4	25.326029	13.779881	16.254766	68.404041
5	19.114166	19.228955	22.455276	62.289232
6	NaN	38.064344	41.487320	43.299401
7	38.064344	NaN	6.793841	81.200755
8	41.487320	6.793841	NaN	84.510798
9	43.299401	81.200755	84.510798	NaN

```

In [14]: distances = pd.DataFrame([dsq.idxmin(axis=0), dsq.min(axis=0), dsq.idxmax(axis=0), dsq.max(axis=0)])
         distances.columns = ['nearest_euclidean_point', 'near_euclidean_dist_val', 'far_euclidean_point', 'far_euclidean_dist_val']
         display(distances)

```

	nearest_euclidean_point	near_euclidean_dist_val	far_euclidean_point	\
0	4.0	0.000000	9.0	
1	3.0	11.257696	8.0	
2	6.0	5.071906	9.0	
3	6.0	4.186459	8.0	
4	0.0	0.000000	9.0	
5	0.0	6.325472	9.0	
6	3.0	4.186459	9.0	



7	8.0	6.793841	9.0
8	7.0	6.793841	9.0
9	1.0	36.085200	8.0

	far_euclidean_dist_val
0	68.404041
1	51.590491
2	48.199803
3	42.907828
4	68.404041
5	62.289232
6	43.299401
7	81.200755
8	84.510798
9	84.510798